

---

---

# **D Perf Patterns**

— S.Rohe Perf Meetup Munich —  
1.2.2017

---

---

# Agenda

- Motivation
- Boundaries
- DLang
- Examples
- Summary

**Motivation**

sociomantic  
from dunnhumby

eBay

facebook



REMEDY


WEKA.io

AdRoll

funkwerk  
aktiengesellschaft

# Motivation (vibe.d - D WebServer)

Requests per second



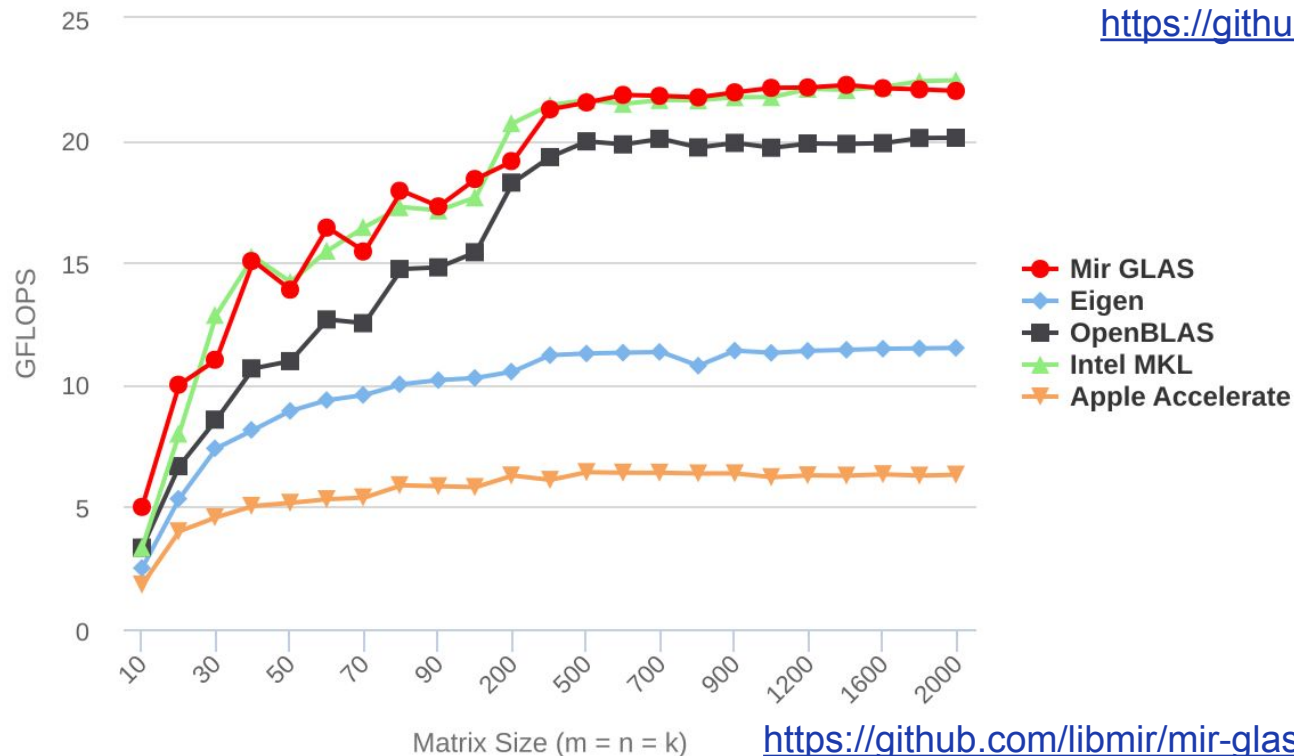
Benchmarks by Михаил Страшун (Dicebot) - <http://j.mp/14SASPX>

<https://github.com/rejectedsoftware/vibe.d>

# Motivation (libmir - Matrix Operations)

Complex single precision

General Matrix-matrix Multiplication



<https://github.com/libmir/mir-cpuid>

<https://github.com/libmir/mir-glas>

# Boundaries

just single core

Performance difficult to define (run time performance, memory consumption, parallel performance, ...)

not particular D specific

no proper measurements for following examples

# DLang

**D** is a systems programming language with C-like syntax and static typing. It combines efficiency, control and modeling power with safety and programmer productivity. [dlang.org]

since 2001 D1; since 2011 D2

designed by C++ Compiler Guru and C++ Template Guru

fast compilation time, static if / static foreach

<http://tour.dlang.org/>



# How to use Prime Numbers?

program needs prime numbers

- compute them so fast, that caching makes no sense
- cache them on first calculation
- read in a precalculated prime.txt file
- compute them during compilation



# DLang - CT-Primes

Example: primes.d

```
#!/usr/bin/env rdmd
import std.algorithm, std.range, std.stdio;


uint[] primes(uint max) {
    return iota(1, max).filter!isPrime.array;
}

void main() {
    enum a = primes(10000);
    stdout.writeln("compile time primes %s", a);
}
```

```
// Algorithm from
https://en.wikipedia.org/wiki/Primality\_test
bool isPrime(uint n) {
    if (n <= 3) return n > 1;
    if (n % 2 == 0 || n % 3 == 0) return false;
    uint i = 5;
    while (i*i <= n) {
        if (n % i == 0 || n % (i + 2) == 0) return false;
        i += 6;
    }
    return true;
}

unittest {
    assert(2.isPrime);
    assert(11.isPrime);
    assert(!12.isPrime);
    assert(97.isPrime);
}
```

# Motivation (regex)



<https://github.com/DmitryOlshansky/FReD>

# Examples - Regular Expressions

```
#!/usr/bin/env rdmd
```

Example: regex.d

```
import std.regex, std.stdio;
```

```
void main() {  
    auto r = ctRegex!`\d\d.\d\d.\d\d\d\d`;   
    auto r2 = regex(`\d\d.\d\d.\d\d\d\d`);  
    string long_text = "this is a long string with some date like 01.02.2003 in it.";   
  
    stdout.writeln("compile time match?: %s", long_text.match(r));  
    stdout.writeln("run time match?: %s", long_text.match(r2));  
}
```

# Pattern 1: Compile Time Evaluation

- regular expressions
- $\frac{1}{3}$ ,  $\frac{1}{2}$ ,  $\text{sqrt}(2)$ , simple calculations
- cpuid
  - array length based on cache sizes
- static for - loop unrolling
- schema validation DSL
- define own grammars / DSLs: <https://github.com/PhilippeSigaud/Pegged>
- ...
- “calculate” algorithms at compile time

# Sort-n

- special sort for a fixed set of members
- sort requires partitioning datasets into smaller partitions
- for large partitions this makes sense
- for smaller partitions:
  - splitting is not efficient anymore
  - do the  $n$  comparisons and swaps by hand is faster

Example: `sortn.d`

## Sort-3

```
static if (n == 3)
{
    s.conditionalSwap!(less, Range,
        0,1,
        1,2,
        0,1);
}
```

## Sort-4

```
static if (n == 4)
{
    s.conditionalSwap!(less, Range,
        0,1, 2,3, // 2 in parallel
        0,2, 1,3, // 2 in parallel
        1,2);
}
```

# Memory is slow

## Size and Speeds

- L1: 3-4 cycles, 32 kb
  - instruction
  - data
- L2: 8-14 cycles, 256kb
- L3: tens of cycles, few MB, often shared
- Cache line: 64 bytes

# Struct Padding

```
struct S {  
    bool f1;  
    uint f2;  
    bool f3;  
}
```



12 bytes, 6 wasted



# Struct Padding

```
struct S {  
    uint f2;  
    bool f1;  
    bool f3;  
}
```



8 bytes, 2 wasted

# Struct Packing

Example: packing.d

```
import std.bitmanip;
```

```
struct S {  
    mixin(bitfield!(  
        uint, "f1", 30,  
        bool, "f2", 1,  
        bool, "f3", 1,  
    ));  
}
```



4 bytes, 0 wasted

- f1 is now 30 bits instead of 32 bits
  - Now about 1B max
- Fields aren't atomic anymore
- bitfield does all the magic

# Class/Struct Inheritance

- Classes are on the heap, require pointer
- Structs are on stack, no pointer
- Structs cannot be inherited, normally


Example: struct.d

## Pattern 2: Reduce Data

- less data, more efficient usage of the cache line
- less abstractions, less data

# Partition I

- Idea of a Sentinel for find
- with sentinel does not need to check if end of range reached
- find with sentinel requires:
  - random access range
  - moving elements

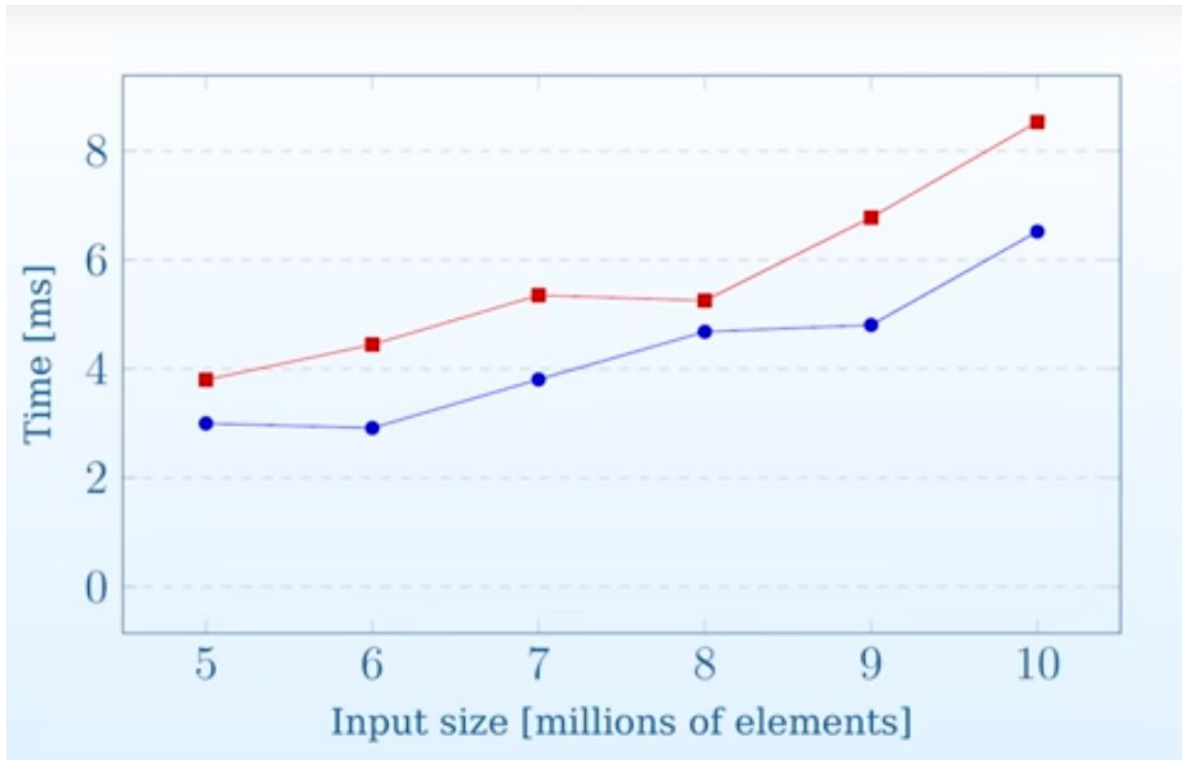


# Partition II - Design By Introspection

## static if to the rescue

```
Range find(R, E)(R r, E e)
if (isInputRange!R && is(typeof(r.front) == e) : bool))
    static if (isRandomAccessRange!R && hasSlicing!R) {
        static if (is(typeof(
            () nothrow { r[0] = r[0]; }
            ))) {
            ... sentinel implementation ...
        } else {
            ... indexed implementation ...
        }
    } else {
        ... conservative implementation ...
    }
}
```

# Partition III - Find with Sentinel




# Partition IV

- partition could be also done via sentinel
- (quick)sort performance relies heavily on partition
- Algorithm:
  - Plant sentinels at both ends
  - Create a “vacancy” in the range
  - Break swapAt into two assignments
  - An assignment fills a vacancy and leaves another one



## Partition V - with sentinel (blue)



# Pattern 3: Design By Introspection

- every function could provide multiple implementations
  - a conservative one
  - an optimized one
- implementation chosen at compile time based on the properties of a parameter types

# Performance Patterns

- Doing as little work as possible
  - lazy log.warn
- Doing as much as possible at Compile Time
  - sqrt(2) at compile time, regex at compile time, design by introspection
- Optimize to the current CPU; cache line sizes, instructions, ...
  - preload the amount of data that fits your cache sizes
  - Vectorization - unfolding for loops based on cache sizes
- Reduce data dependency
  - sortn
- Reduce size of data
  - int8 instead of int64, bitpacking
- Reduce abstractions
  - struct inheritance



# For D Lovers

## Next Meetup

**DLang and the Cloud - 14.02.2017 @**

**celonis**

- **DLang and the Cloud**
- **D-Compiler usage within the Cloud**
- **Heroku with D**

# References

**Amaury Sechet - Bit Packing like a Madman**

[https://www.youtube.com/watch?v=95O\\_y9fu6qk](https://www.youtube.com/watch?v=95O_y9fu6qk)

**Alexandrescu - Writing Fast Code - how to optimize atoi**

Part 1: <https://www.youtube.com/watch?v=vrfYLIR8X8k>

Part 2: <https://www.youtube.com/watch?v=9tvbz8CSl8M>

**Alexandrescu - Fastware - ACCU 2016 Keynote**

<https://www.youtube.com/watch?v=AxnotgLql0k>